

notebook

May 8, 2022

0.1 Imports

```
[ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import ipaddress

from IPython.display import display
from sklearn.covariance import EllipticEnvelope
from sklearn.ensemble import IsolationForest
from sklearn.neighbors import LocalOutlierFactor
from sklearn.metrics import precision_score, recall_score, accuracy_score,
    ↪ confusion_matrix, average_precision_score, ConfusionMatrixDisplay,
    ↪ roc_curve, auc, PrecisionRecallDisplay

rng = np.random.RandomState(42)
```

0.2 Data loading

```
[ ]: col_names = ['end_epoch', 'duration', 'src_ip', 'dst_ip', 'src_port',
    ↪ 'dst_port', 'protocol', 'flags', 'status', 'service_type', 'packets',
    ↪ 'bytes', 'attack']

def load_dataframe(name):
    return pd.read_csv(f"data/training/{name}_flows_cut.csv", header = None,
    ↪ names = col_names)

def load_test_dataframe(name):
    return pd.read_csv(f"data/test/{name}_flows_cut.csv", header = None, names
    ↪ = col_names)
```

```
[ ]: # Load training datasets
background = load_dataframe("background")
blacklist = load_dataframe("blacklist")
sshscan = load_dataframe("sshscan")
spam = load_dataframe("spam")
```

```

# Load tests datasets
test_background = load_test_dataframe("background")
test_blacklist = load_test_dataframe("blacklist")
test_sshscan = load_test_dataframe("sshscan")
test_botnet = load_test_dataframe("botnet")
test_scan11 = load_test_dataframe("scan11")
test_scan44 = load_test_dataframe("scan44")
test_spam = load_test_dataframe("spam")
test_dos = load_test_dataframe("dos")

# We regroup all attacks of each dataset in lists to have a cleaner code latter
attacks_dataframes = [blacklist, sshscan, spam]
test_attacks_dataframes = [test_blacklist, test_sshscan, test_botnet,
↳test_scan11, test_scan44, test_spam, test_dos]

```

```

[ ]: def get_dataframes(outlier_ratio, num_samples):
    """
    This function allow us to get two proper datasets of training and tests
    ↳with the given size and outlier ratio.
    """
    # Set offsets of each dataset depending on poor sized sub-datasets
    train_offset = len(spam) % num_samples * outlier_ratio
    test_offset = len(test_sshscan) % num_samples * outlier_ratio

    # Get the number of attack sample our datasets will contains
    nb_train_attack = int((num_samples * outlier_ratio - train_offset) //
↳(len(attacks_dataframes) - 1))
    nb_test_attack = int((num_samples * outlier_ratio - test_offset) //
↳(len(test_attacks_dataframes) - 1))

    # Get the correct number of sample for each attack sub-dataset
    train_dataframes = [df.head(nb_train_attack) for df in attacks_dataframes]
    test_dataframes = [df.head(nb_test_attack) for df in
↳test_attacks_dataframes]

    # The remaining of space allowed is fill with background flows for each
    ↳dataset
    train_dataframes.append(background.head(num_samples - sum(len(df) for df in
↳train_dataframes)))
    test_dataframes.append(test_background.head(num_samples - sum(len(df) for
↳df in test_dataframes)))

    # Concat each lists to get our dataframes
    train_df = pd.concat(train_dataframes, ignore_index=True)
    test_df = pd.concat(test_dataframes, ignore_index=True)

```

```
return train_df, test_df
```

```
[ ]: outlier_ratio = 0.1
train_df, test_df = get_dataframes(outlier_ratio, 10000)
train_df.head()
```

```
[ ]:      end_epoch  duration      src_ip      dst_ip  src_port  \
0  2016-04-18 00:02:16    0.996  42.219.158.242    60.56.180.24    33421
1  2016-04-18 00:02:16    0.852  42.219.158.242    60.56.180.24    35297
2  2016-04-18 00:02:17    0.936    60.56.180.24  42.219.158.242      80
3  2016-04-18 00:02:17    0.804    60.56.180.24  42.219.158.242    443
4  2016-04-18 00:02:30    0.000  42.219.152.242  88.205.102.190   38531

      dst_port  protocol  flags  status  service_type  packets  bytes  attack
0         80      TCP  .AP.SF      0          0         6    437  blacklist
1        443      TCP  .AP.SF      0          0        52   3030  blacklist
2       33421      TCP  .AP.SF      0          0         4    565  blacklist
3       35297      TCP  .AP.SF      0          0        72  86267  blacklist
4         25      TCP  ...S.      0          0         1    60  blacklist
```

0.3 Data analysis

As described in the scientific paper, we know that our ‘background’ data may contain unlabeled attacks.

Because we can't know for sure which exchanges are attacks we'll consider them as genuine exchanges for now.

Same goes for blacklist flows, we can't know for sure if they are attacks we'll consider them as fraudulent exchanges for now.

0.3.1 Training dataset observations

```
[ ]: print(f'The training dataset size is: {len(train_df)}\n')

print(f'The number of unknown attacks / background exchanges is:␣
↪{len(train_df[train_df.attack == "background"])}')
print(f'The number of known attacks is: {len(train_df[train_df.attack !=␣
↪"background"])}\n')

print(f'The types of attack are: {train_df[train_df.attack != "background"].
↪attack.unique()}')
for attack in train_df[train_df.attack != "background"].attack.unique():
    print(f'The number of {attack} is: {len(train_df[train_df.attack ==␣
↪attack])}')
print()

print(f'Number of source IP used in known attacks: {len(train_df[train_df.
↪attack != "background"].src_ip.unique())}')

```

```

print(f'Number of source IP used in othe exchanges: {len(train_df[train_df.
↪attack == "background"].src_ip.unique()})')
print(f'Number of destination IP used in known attacks: {len(train_df[train_df.
↪attack != "background"].dst_ip.unique()})')
print(f'Number of destination IP used in other exchanges:␣
↪{len(train_df[train_df.attack == "background"].dst_ip.unique()})\n')

print(f'Number of source port used in known attacks: {len(train_df[train_df.
↪attack != "background"].src_port.unique()})')
print(f'Number of source port used in othe exchanges: {len(train_df[train_df.
↪attack == "background"].src_port.unique()})')
print(f'Number of destination port used in known attacks:␣
↪{len(train_df[train_df.attack != "background"].dst_port.unique()})')
print(f'Number of destination port used in othe exchanges:␣
↪{len(train_df[train_df.attack == "background"].dst_port.unique()})\n')

print(f'Protocols used during known attacks: {train_df[train_df.attack !=␣
↪"background"].protocol.unique()})')
print(f'Protocols used during other exchanges: {train_df[train_df.attack ==␣
↪"background"].protocol.unique()})\n')

```

The training dataset size is: 10000

The number of unknown attacks / background exchanges is: 8965

The number of known attacks is: 1035

The types of attack are: ['blacklist' 'anomaly-sshscan' 'anomaly-spam']

The number of blacklist is: 497

The number of anomaly-sshscan is: 497

The number of anomaly-spam is: 41

Number of source IP used in known attacks: 251

Number of source IP used in othe exchanges: 2018

Number of destination IP used in known attacks: 227

Number of destination IP used in other exchanges: 1537

Number of source port used in known attacks: 451

Number of source port used in othe exchanges: 4363

Number of destination port used in known attacks: 495

Number of destination port used in othe exchanges: 3939

Protocols used during known attacks: ['TCP' 'UDP' 'ICMP']

Protocols used during other exchanges: ['TCP' 'UDP' 'GRE' 'ESP' 'ICMP' 'IPv6']

```
[ ]: print(f'Flags used during known attacks: {train_df[train_df.attack != "background"]["flags"].unique()}')
      print(f'Flags used during other exchanges: {train_df[train_df.attack == "background"]["flags"].unique()}\n')

      print(f'Types of service used during known attacks: {train_df[train_df.attack != "background"].service_type.unique()}')
      print(f'Types of service used during other exchanges: {train_df[train_df.attack == "background"].service_type.unique()}\n')

      print(f'Number of different status in the dataset: {len(train_df.status.unique())}\n')

      print('Details of quantitative features of known attacks:')
      display(train_df[train_df.attack != "background"][['duration', 'packets', 'bytes']].describe())
      print()
      print('Details of quantitative features of other exchanges:')
      display(train_df[train_df.attack == "background"][['duration', 'packets', 'bytes']].describe())
```

```
Flags used during known attacks: ['.AP.SF' '...S.' '.A.R..' '.AP.S.' '.A...F'
'.A...' '.A..SF' '.APRSF'
'.A..S.' '.AP..F' '.AP...' '.APR.F' '...R..' '.APRS.' '...RS.']
Flags used during other exchanges: ['.A...' '.A..S.' '.AP.S.' '.AP...' '...S.'
'.A...F' '.A.R.F' '.APR..'
'.A.R..' '.AP..F' '...R..' '.AP.SF' '.A..SF' '.APRS.' '.APRSF' '.A.RSF'
'.APR.F' '.A.RS.']
```

```
Types of service used during known attacks: [ 0 40 72  8]
Types of service used during other exchanges: [  0  64  40  72   2   8 192  26
24 200  80  42 184  16 104  20  75]
```

```
Number of different status in the dataset: 1
```

```
Details of quantitative features of known attacks:
```

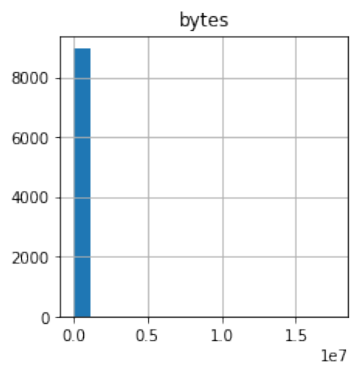
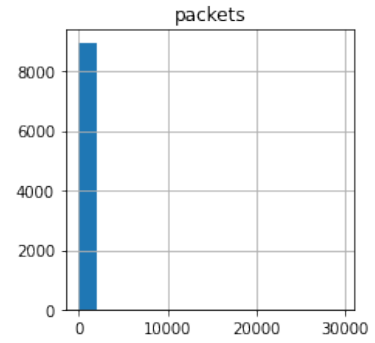
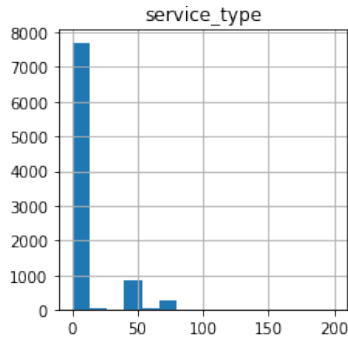

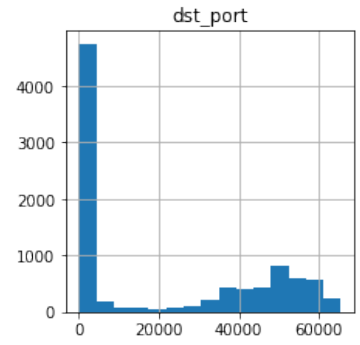


	duration	packets	bytes
count	1035.000000	1035.000000	1.035000e+03
mean	4.668881	24.860870	1.485774e+04
std	5.745437	63.051175	7.601059e+04
min	0.000000	1.000000	4.000000e+01
25%	0.888000	5.000000	5.650000e+02
50%	3.540000	12.000000	1.373000e+03
75%	5.040000	15.000000	1.954000e+03
max	54.712000	782.000000	1.118209e+06

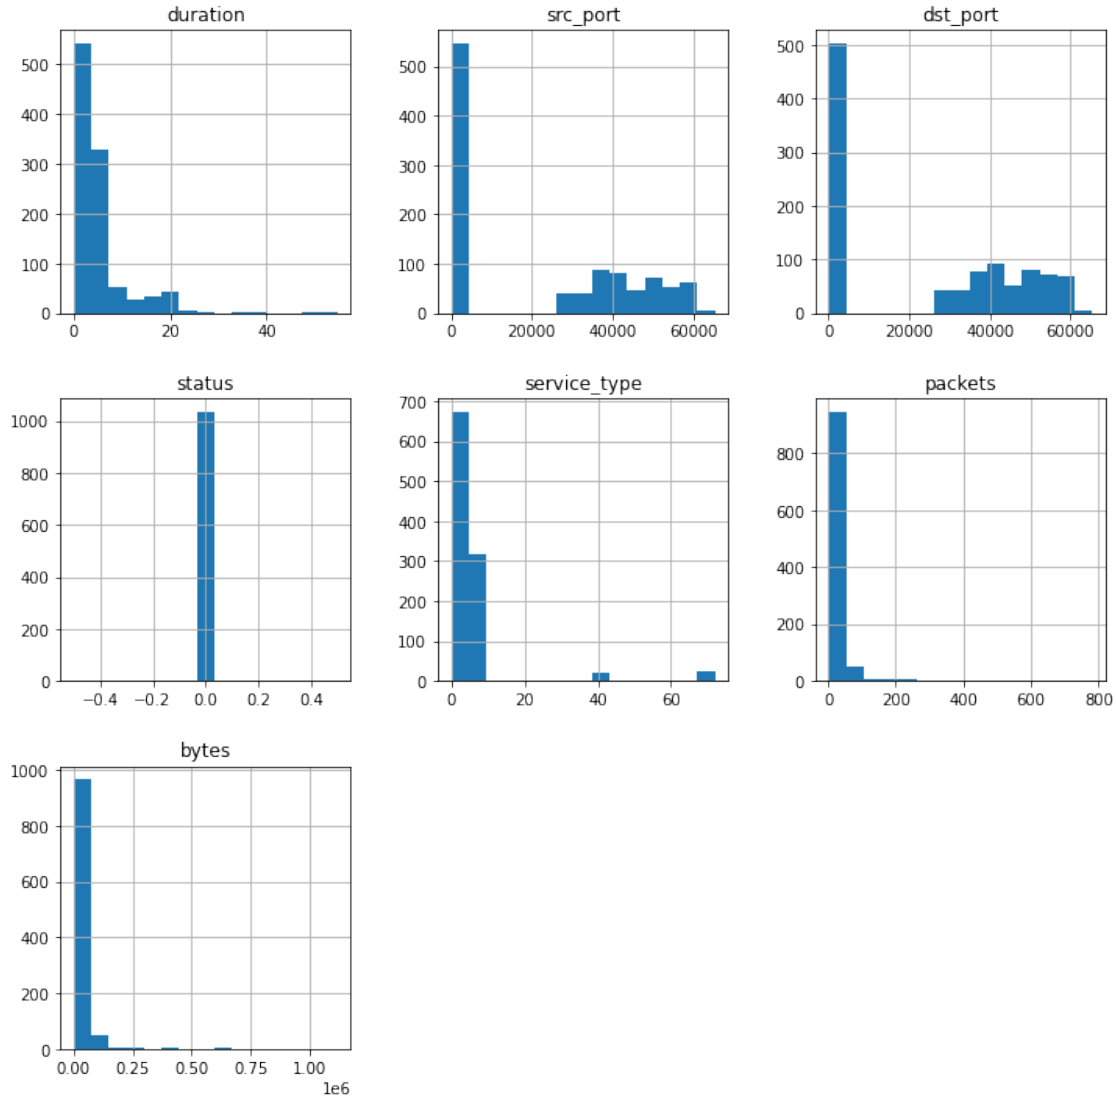
Details of quantitative features of other exchanges:

	duration	packets	bytes
count	8965.000000	8965.000000	8.965000e+03
mean	1.109757	15.008589	7.854498e+03
std	5.411496	374.938073	2.400408e+05
min	0.000000	1.000000	3.100000e+01
25%	0.000000	1.000000	7.600000e+01
50%	0.000000	1.000000	1.310000e+02
75%	0.324000	6.000000	7.280000e+02
max	194.796000	29544.000000	1.765682e+07

```
[ ]: train_df[train_df.attack == "background"].hist(bins=15, figsize=(12, 12))
train_df[train_df.attack != "background"].hist(bins=15, figsize=(12, 12))
```

```
[ ]: array([[<AxesSubplot:title={'center': 'duration'}>,
            <AxesSubplot:title={'center': 'src_port'}>,
            <AxesSubplot:title={'center': 'dst_port'}>],
            [<AxesSubplot:title={'center': 'status'}>,
            <AxesSubplot:title={'center': 'service_type'}>,
            <AxesSubplot:title={'center': 'packets'}>],
            [<AxesSubplot:title={'center': 'bytes'}>, <AxesSubplot:>,
            <AxesSubplot:>]], dtype=object)
```





0.3.2 Test Dataset observations

```
[ ]: print(f'The test dataset size is: {len(test_df)}\n')

print(f'The number of unknown attacks / background exchanges is:␣
↪{len(test_df[test_df.attack == "background"])}')
print(f'The number of known attacks is: {len(test_df[test_df.attack !=␣
↪"background"])}\n')

print(f'The types of attack are: {test_df[test_df.attack != "background"].
↪attack.unique()}')
for attack in test_df[test_df.attack != "background"].attack.unique():
```



```

    print(f'The number of {attack} is: {len(test_df[test_df.attack ==
↪attack])}')
print()

print(f'Number of source IP used in known attacks: {len(test_df[test_df.attack !=
↪ "background"].src_ip.unique()})')
print(f'Number of source IP used in othe exchanges: {len(test_df[test_df.attack ==
↪ "background"].src_ip.unique()})')
print(f'Number of destination IP used in known attacks: {len(test_df[test_df.
↪attack != "background"].dst_ip.unique()})')
print(f'Number of destination IP used in other exchanges: {len(test_df[test_df.
↪attack == "background"].dst_ip.unique()})\n')

print(f'Number of source port used in known attacks: {len(test_df[test_df.
↪attack != "background"].src_port.unique()})')
print(f'Number of source port used in othe exchanges: {len(test_df[test_df.
↪attack == "background"].src_port.unique()})')
print(f'Number of destination port used in known attacks: {len(test_df[test_df.
↪attack != "background"].dst_port.unique()})')
print(f'Number of destination port used in othe exchanges: {len(test_df[test_df.
↪attack == "background"].dst_port.unique()})\n')

print(f'Protocols used during known attacks: {test_df[test_df.attack !=
↪ "background"].protocol.unique()})')
print(f'Protocols used during other exchanges: {test_df[test_df.attack ==
↪ "background"].protocol.unique()})\n')

```

The test dataset size is: 10000

The number of unknown attacks / background exchanges is: 8999

The number of known attacks is: 1001

The types of attack are: ['blacklist' 'anomaly-sshscan' 'nerisbotnet' 'scan11' 'scan44'

'anomaly-spam' 'dos']

The number of blacklist is: 166

The number of anomaly-sshscan is: 5

The number of nerisbotnet is: 166

The number of scan11 is: 166

The number of scan44 is: 166

The number of anomaly-spam is: 166

The number of dos is: 166

Number of source IP used in known attacks: 131

Number of source IP used in othe exchanges: 1603

Number of destination IP used in known attacks: 50

Number of destination IP used in other exchanges: 3311

Number of source port used in known attacks: 515
 Number of source port used in othe exchanges: 4455
 Number of destination port used in known attacks: 274
 Number of destination port used in othe exchanges: 3838

Protocols used during known attacks: ['TCP' 'UDP' 'ICMP']
 Protocols used during other exchanges: ['ICMP' 'TCP' 'UDP']

```
[ ]: print(f'Flags used during known attacks: {test_df[test_df.attack != "background"]["flags"].unique()}')
      print(f'Flags used during other exchanges: {test_df[test_df.attack == "background"]["flags"].unique()}\n')

      print(f'Types of service used during known attacks: {test_df[test_df.attack != "background"].service_type.unique()}')
      print(f'Types of service used during other exchanges: {test_df[test_df.attack == "background"].service_type.unique()}\n')

      print(f'Number of different status in the dataset: {len(test_df.status.unique())}\n')

      print('Details of quantitative features of known attacks:')
      display(test_df[test_df.attack != "background"][['duration', 'packets', 'bytes']].describe())
      print()
      print('Details of quantitative features of other exchanges:')
      display(test_df[test_df.attack == "background"][['duration', 'packets', 'bytes']].describe())
```

Flags used during known attacks: ['.APRS.' '.AP.S.' '.APR..' '.AP.SF' '.A..SF' '...S.' '.APRSF' '.A.R..' '.A..' '.A.RSF' '...R..' '.AP..F' '.A..S.' '...RS.']
 Flags used during other exchanges: ['.A..' '.A..S.' '...S.' '.A..F' '.A.R..' '.APR..' '.AP..' '...R..' '.AP.S.' '.A.RS.' '.A..SF' '.AP.SF' '.APR.F' '.APRSF' '.AP..F' '.APRS.' '.A.R.F' '...RS.']

Types of service used during known attacks: [40 0 72 8 74]
 Types of service used during other exchanges: [192 0 200 26 2 64 24 72 40 4 8 20 16]

Number of different status in the dataset: 1

Details of quantitative features of known attacks:

duration	packets	bytes
----------	---------	-------

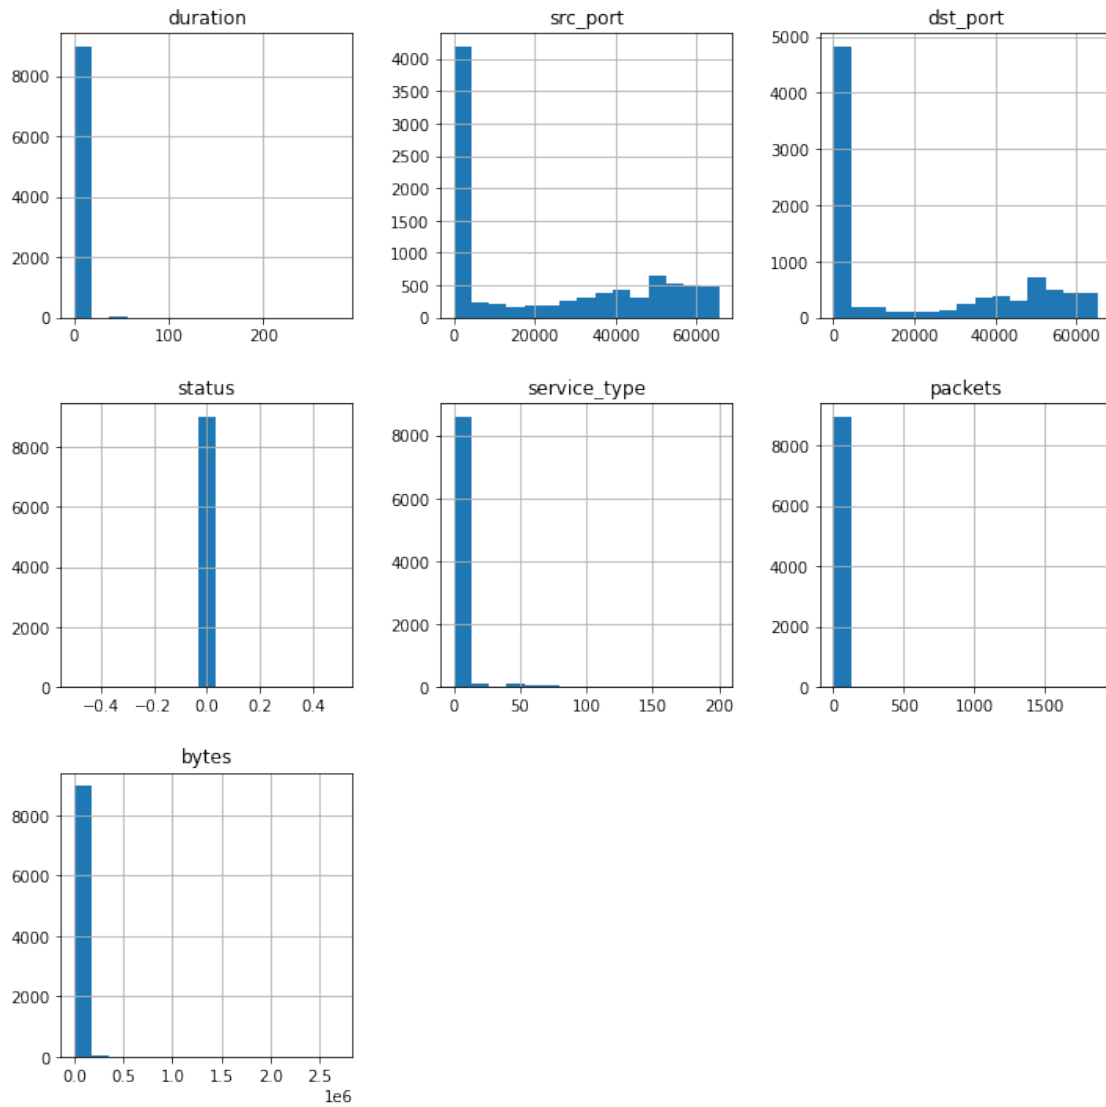
count	1001.000000	1001.000000	1001.000000
mean	0.779888	3.335664	634.646354
std	2.593195	3.860987	2776.405061
min	0.000000	1.000000	28.000000
25%	0.000000	1.000000	44.000000
50%	0.000000	2.000000	197.000000
75%	0.576000	6.000000	441.000000
max	21.060000	27.000000	36132.000000

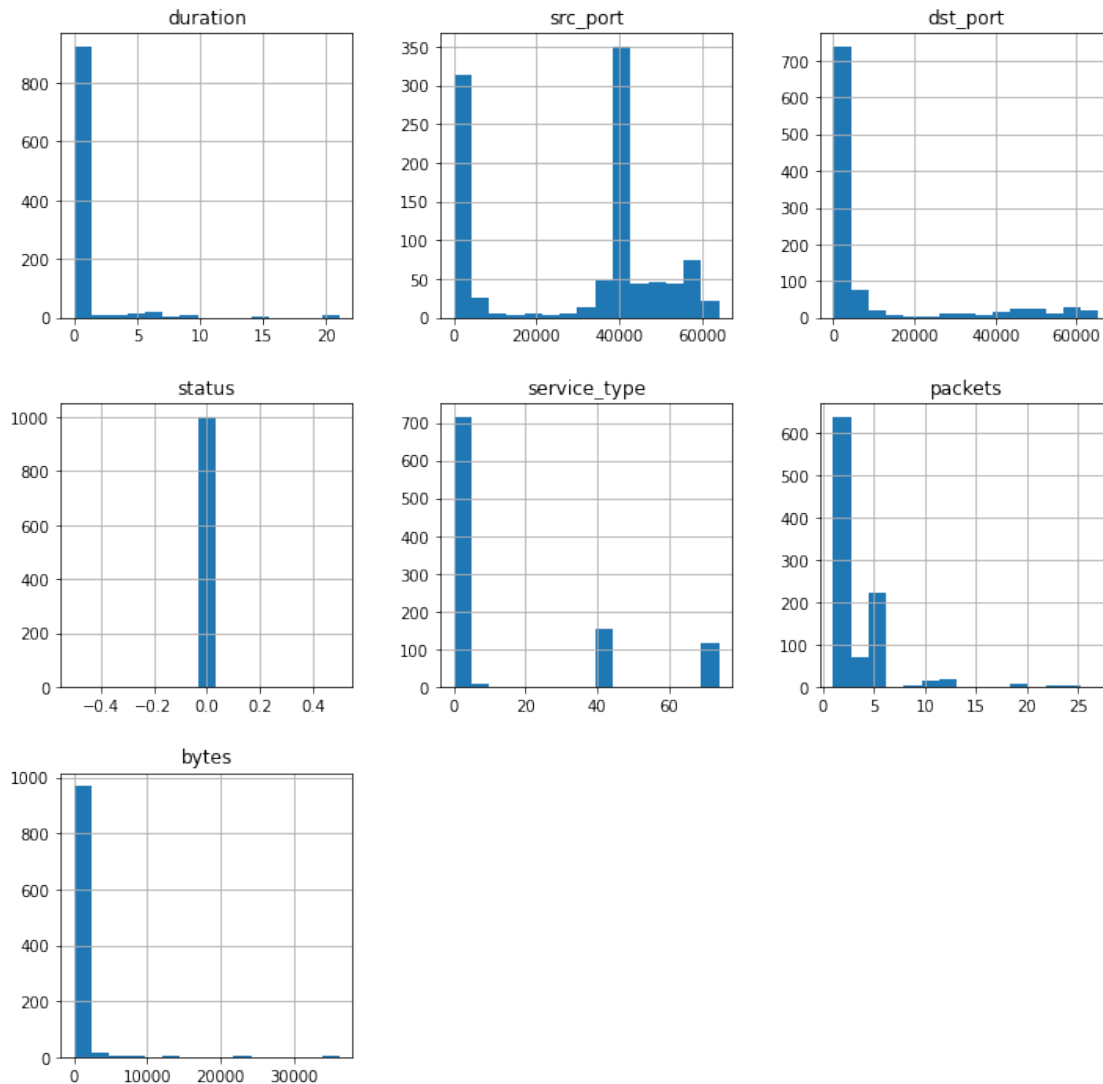
Details of quantitative features of other exchanges:

	duration	packets	bytes
count	8999.000000	8999.000000	8.999000e+03
mean	1.431335	7.270252	4.047749e+03
std	5.636508	31.983570	4.009776e+04
min	0.000000	1.000000	2.800000e+01
25%	0.000000	1.000000	9.600000e+01
50%	0.200000	2.000000	1.640000e+02
75%	1.828000	7.000000	1.081500e+03
max	280.620000	1879.000000	2.703238e+06

```
[ ]: test_df[test_df.attack == "background"].hist(bins=15, figsize=(12, 12))
test_df[test_df.attack != "background"].hist(bins=15, figsize=(12, 12))
```

```
[ ]: array([[<AxesSubplot:title={'center': 'duration'}>,
<AxesSubplot:title={'center': 'src_port'}>,
<AxesSubplot:title={'center': 'dst_port'}>],
[<AxesSubplot:title={'center': 'status'}>,
<AxesSubplot:title={'center': 'service_type'}>,
<AxesSubplot:title={'center': 'packets'}>],
[<AxesSubplot:title={'center': 'bytes'}>, <AxesSubplot:>,
<AxesSubplot:>]], dtype=object)
```





0.4 Data Cleaning

```
[ ]: # We drop 'status' column as we saw that there is only one value accepted all
      ↪ datasets
      # We also drop 'packets' column as it is redondant with the 'bytes' column

train_df.drop(['status'], axis=1, inplace=True)
test_df.drop(['status'], axis=1, inplace=True)

print(f'Number of missing values in the train dataset: {train_df.isnull().
      ↪ values.sum()}')
print(f'Number of missing values in the test dataset: {test_df.isnull().values.
      ↪ sum()}')
```

Number of missing values in the train dataset: 0

Number of missing values in the test dataset: 0

```
[ ]: def encode_dataframes(train_df, test_df):
    """
    Encode the training and testing dataframes to have better performance in
    our algorithms.
    """
    train_df = train_df.drop(['attack'], axis=1)
    test_df = test_df.drop(['attack'], axis=1)

    # Convert timestamp to epoch integer, we dont keep the date values to not
    have biased models later
    train_df.end_epoch = pd.to_datetime(train_df.end_epoch).apply(lambda x: x.
    strftime('%H%M%S')).astype(np.int64)
    test_df.end_epoch = pd.to_datetime(test_df.end_epoch).apply(lambda x: x.
    strftime('%H%M%S')).astype(np.int64)

    # Convert IPs to int, for optimization purpose only
    train_df.src_ip = train_df.src_ip.apply(lambda x: np.int64(ipaddress.
    IPv4Address(x)))
    train_df.dst_ip = train_df.dst_ip.apply(lambda x: np.int64(ipaddress.
    IPv4Address(x)))
    test_df.src_ip = test_df.src_ip.apply(lambda x: np.int64(ipaddress.
    IPv4Address(x)))
    test_df.dst_ip = test_df.dst_ip.apply(lambda x: np.int64(ipaddress.
    IPv4Address(x)))

    encoded_train_df = pd.get_dummies(train_df)
    encoded_test_df = pd.get_dummies(test_df)

    # Add missing columns to each dataset for consistency
    for column in encoded_train_df.columns:
        if column not in encoded_test_df.columns:
            encoded_test_df[column] = 0

    for column in encoded_test_df.columns:
        if column not in encoded_train_df.columns:
            encoded_train_df[column] = 0

    # Reorder columns order, also for consistency
    encoded_test_df = encoded_test_df[encoded_train_df.columns]

    return encoded_train_df, encoded_test_df

encoded_train_df, encoded_test_df = encode_dataframes(train_df, test_df)
encoded_train_df.head()
```

```
[ ]:  end_epoch  duration      src_ip      dst_ip  src_port  dst_port  \
0      216      0.996   719036146  1010349080    33421      80
1      216      0.852   719036146  1010349080    35297     443
2      217      0.936  1010349080   719036146      80    33421
3      217      0.804  1010349080   719036146     443    35297
4      230      0.000   719034610  1489856190    38531     25

      service_type  packets  bytes  protocol_ESP  ...  flags_.A.RS.  \
0              0        6    437              0  ...              0
1              0       52   3030              0  ...              0
2              0        4    565              0  ...              0
3              0       72  86267              0  ...              0
4              0        1    60              0  ...              0

      flags_.A.RSF  flags_.AP...  flags_.AP..F  flags_.AP.S.  flags_.AP.SF  \
0              0              0              0              0              1
1              0              0              0              0              1
2              0              0              0              0              1
3              0              0              0              0              1
4              0              0              0              0              0

      flags_.APR..  flags_.APR.F  flags_.APRS.  flags_.APRSF
0              0              0              0              0
1              0              0              0              0
2              0              0              0              0
3              0              0              0              0
4              0              0              0              0
```

[5 rows x 34 columns]

0.5 Dataset Visualisation

```
[ ]: def show_flow_exchanges_3d(df, df_attack, x_axis, y_axis, z_axis, title):
      fig = plt.figure(figsize=(10,10))
      ax = fig.add_subplot(111, projection='3d')

      df_attacks = df[df_attack != "background"]
      xa = df_attacks[x_axis]
      ya = df_attacks[y_axis]
      za = df_attacks[z_axis]

      ax.scatter3D(np.log10(xa + 0.1), np.log10(ya + 0.1), np.log10(za + 0.1),
        ↪label="attack", color="r", marker=".", s=15)

      df_genuines = df[df_attack == "background"]
      xg = df_genuines[x_axis]
      yg = df_genuines[y_axis]
```

```

zg = df_genuines[z_axis]

ax.scatter3D(np.log10(xg + 0.1), np.log10(yg + 0.1), np.log10(zg + 0.1),
↳label="background", color="b", marker=".", s=15)

ax.set_title(title, size=16)
ax.set_xlabel('log$_{10}$ ( ' + x_axis + ' )', size=14)
ax.set_ylabel('log$_{10}$ ( ' + y_axis + ' )', size=14)
ax.set_zlabel('log$_{10}$ ( ' + z_axis + ' )', size=14)

plt.legend()


```

```


[ ]: # End epoch differs a lot between attacks and normal flows
show_flow_exchanges_3d(encoded_train_df, train_df["attack"], 'end_epoch',
↳'bytes', 'duration', 'Representation of all exchanges in the training
↳dataset')
show_flow_exchanges_3d(encoded_test_df, test_df["attack"], 'end_epoch',
↳'bytes', 'duration', 'Representation of all exchanges in the testing
↳dataset')

```


Representation of all exchanges in the training dataset




Representation of all exchanges in the testing dataset




```
[ ]: # Packets differs a bit between attacks and normal flows
show_flow_exchanges_3d(encoded_train_df, train_df["attack"], 'packets', 'bytes', 'duration', 'Representation of all exchanges in the training dataset')
show_flow_exchanges_3d(encoded_test_df, test_df["attack"], 'packets', 'bytes', 'duration', 'Representation of all exchanges in the testing dataset')
```

Representation of all exchanges in the training dataset




Representation of all exchanges in the testing dataset




```
[ ]: # Attacks seems to be focused on a reduced set of source / destination ports.
show_flow_exchanges_3d(encoded_train_df, train_df["attack"], 'dst_port', 'src_port', 'duration', 'Representation of all exchanges in the training dataset')
show_flow_exchanges_3d(encoded_test_df, test_df["attack"], 'dst_port', 'src_port', 'duration', 'Representation of all exchanges in the testing dataset')
```

Representation of all exchanges in the training dataset



Representation of all exchanges in the testing dataset



```
[ ]: # Matrix de correlation
corr = train_df.corr()
corr.style.background_gradient(cmap="coolwarm")
```

```
[ ]: <pandas.io.formats.style.Styler at 0x7f21e1db1a90>
```

0.6 Outlier detection

```
[ ]: def display_metrics(y_true, y_pred):
    fig, axes = plt.subplots(1, 3, figsize=(20, 5))

    # Display the confusion matrix
    cm = confusion_matrix(y_true, y_pred)
```

```

tn, fp, fn, tp = cm.ravel()
ConfusionMatrixDisplay(confusion_matrix=cm).plot(ax=axes[0])

# Display classifier metrics
print(f'AUPRC: {round(average_precision_score(y_true, y_pred) * 100, 2) }%')
print(f'Precision: {round(precision_score(y_true, y_pred) * 100, 2) }%')
print(f'Recall: {round(recall_score(y_true, y_pred) * 100, 2) }%')
print(f'True Negative Rate: {round( tn * 100 / (tn + fp), 2) }%')
print(f'Accuracy: {round(accuracy_score(y_true, y_pred) * 100, 2) }%\n')

# Plot ROC curve
fpr, tpr, _ = roc_curve(y_true, y_pred)
roc_auc = round(auc(fpr, tpr), 2)
axes[1].plot([0, 1], [0, 1], color="navy", linestyle="--")
axes[1].plot(fpr, tpr, color='darkorange', label=f'Roc curve (area = {
    ↪roc_auc})')

axes[1].set_xlim([0.0, 1.0])
axes[1].set_ylim([0.0, 1.5])
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('Receiver operating characteristic curve')
axes[1].legend(loc="lower right")

# Plot Precision-Recall curve
PrecisionRecallDisplay.from_predictions(y_true, y_pred, ax=axes[2])
axes[2].set_title('Precision-Recall Curve')

```

Isolation Forests

```

[ ]: classifier = IsolationForest(contamination=outlier_ratio, n_estimators=1000,
    ↪random_state=rng).fit(encoded_train_df.values)
y_pred = classifier.predict(encoded_test_df.values)

df_outliers = encoded_test_df.copy()
df_outliers["attack"] = test_df.attack != "background"
df_outliers['if_outliers'] = [i == -1 for i in y_pred]
df_outliers.head()

```

```

[ ]:
end_epoch  duration      src_ip      dst_ip  src_port  dst_port  \
0         1103      8.152  1167153979  719035603    38490    443
1         1107      4.372  1167153979  719035603    40516    443
2         1132      0.000  1167153979  719035603    40516    443
3         1136      4.464  1167153979  719035603    51560    443
4         1149      0.000  1167153979  719035603    51560    443

service_type  packets  bytes  protocol_ESP  ...  flags_.AP...  \

```

0	40	13	1620	0	...	0
1	40	11	1502	0	...	0
2	40	2	135	0	...	0
3	40	10	1439	0	...	0
4	40	2	135	0	...	0


	flags_.AP..F	flags_.AP.S.	flags_.AP.SF	flags_.APR..	flags_.APR.F	\
0	0	0	0	0	0	
1	0	1	0	0	0	
2	0	0	0	1	0	
3	0	1	0	0	0	
4	0	0	0	1	0	

	flags_.APRS.	flags_.APRSF	attack	if_outliers
0	1	0	True	True
1	0	0	True	False
2	0	0	True	False
3	0	0	True	True
4	0	0	True	True

[5 rows x 36 columns]

```
[ ]: display_metrics(df_outliers.attack, df_outliers.if_outliers)
```

AUPRC: 40.99%
Precision: 47.79%
Recall: 82.02%
True Negative Rate: 90.03%
Accuracy: 89.23%



Elliptic Envelope

```
[ ]: robustCovariance = EllipticEnvelope(contamination=outlier_ratio,
    ↪random_state=rng).fit(encoded_train_df.values)
```



```
rc_outliers = robustCovariance.predict(encoded_test_df.values)

df_outliers['rc_outliers'] = [i == -1 for i in rc_outliers]
df_outliers.head()
```

```
/home/leiyks/.local/lib/python3.8/site-
packages/sklearn/covariance/_robust_covariance.py:738: UserWarning: The
covariance matrix associated to your dataset is not full rank
warnings.warn(
```

```
[ ]:  end_epoch  duration      src_ip      dst_ip  src_port  dst_port  \
0      1103      8.152  1167153979  719035603      38490      443
1      1107      4.372  1167153979  719035603      40516      443
2      1132      0.000  1167153979  719035603      40516      443
3      1136      4.464  1167153979  719035603      51560      443
4      1149      0.000  1167153979  719035603      51560      443

      service_type  packets  bytes  protocol_ESP  ...  flags_.AP..F  \
0              40        13   1620              0  ...              0
1              40        11   1502              0  ...              0
2              40         2    135              0  ...              0
3              40        10   1439              0  ...              0
4              40         2    135              0  ...              0


      flags_.AP.S.  flags_.AP.SF  flags_.APR..  flags_.APR.F  flags_.APRS.  \
0              0              0              0              0              1
1              1              0              0              0              0
2              0              0              1              0              0
3              1              0              0              0              0
4              0              0              1              0              0

      flags_.APRSF  attack  if_outliers  rc_outliers
0              0    True        True        False
1              0    True        False        False
2              0    True        False        False
3              0    True        True        False
4              0    True        True        False
```

[5 rows x 37 columns]

```
[ ]: display_metrics(df_outliers.attack, df_outliers.rc_outliers)
```

```
AUPRC: 39.71%
Precision: 42.08%
Recall: 92.61%
True Negative Rate: 85.82%
Accuracy: 86.5%
```



Local Outlier Factor

```
[ ]: localOutlierFactor = LocalOutlierFactor(contamination=outlier_ratio,
      ↪n_neighbors=1000, n_jobs=-1)

lof_outliers = localOutlierFactor.fit_predict(encoded_test_df.values)

df_outliers['lof_outliers'] = [i == -1 for i in lof_outliers]
df_outliers.head()
```

```
[ ]:   end_epoch  duration      src_ip      dst_ip  src_port  dst_port  \
0      1103      8.152  1167153979  719035603    38490     443
1      1107      4.372  1167153979  719035603    40516     443
2      1132      0.000  1167153979  719035603    40516     443
3      1136      4.464  1167153979  719035603    51560     443
4      1149      0.000  1167153979  719035603    51560     443

   service_type  packets  bytes  protocol_ESP  ...  flags_.AP.S.  \
0             40       13   1620             0  ...           0
1             40       11   1502             0  ...           1
2             40        2    135             0  ...           0
3             40       10   1439             0  ...           1
4             40        2    135             0  ...           0

   flags_.AP.SF  flags_.APR..  flags_.APR.F  flags_.APRS.  flags_.APRSF  \
0              0           0           0           1           0
1              0           0           0           0           0
2              0           1           0           0           0
3              0           0           0           0           0
4              0           1           0           0           0


   attack  if_outliers  rc_outliers  lof_outliers
0     True         True         False         True
1     True        False         False         True
2     True        False         False         True
```

3	True	True	False	True
4	True	True	False	True

[5 rows x 38 columns]

```
[ ]: display_metrics(df_outliers.attack, df_outliers.lof_outliers)
```

AUPRC: 14.08%
Precision: 25.8%
Recall: 25.77%
True Negative Rate: 91.75%
Accuracy: 85.15%



0.7 Algorithm comparison

0.7.1 Number of attacks detected as outliers by all algorithms

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['if_outliers']_
    ↳ == True) & (df_outliers['lof_outliers'] == True) &_
    ↳ (df_outliers['rc_outliers'] == True)])
```

[]: 202

0.7.2 Number of attacks detected as outliers by LOF and RC

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['lof_outliers']_
    ↳ == True) & (df_outliers['rc_outliers'] == True)])
```

[]: 203

0.7.3 Number of attacks detected as outliers by RC and IF

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['if_outliers']_
↪ == True) & (df_outliers['rc_outliers'] == True)])
```

```
[ ]: 770
```

0.7.4 Number of attacks detected as outliers by LOF and IF

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['if_outliers']_
↪ == True) & (df_outliers['lof_outliers'] == True)])
```

```
[ ]: 244
```

0.7.5 Number of attacks detected as outliers by LOF

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['lof_outliers']_
↪ == True)])
```

```
[ ]: 258
```

0.7.6 Number of attacks detected as outliers by IF

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['if_outliers']_
↪ == True)])
```

```
[ ]: 821
```

0.7.7 Number of attacks detected as outliers by RC

```
[ ]: len(df_outliers[(df_outliers['attack'] == True) & (df_outliers['rc_outliers']_
↪ == True)])
```

```
[ ]: 927
```

0.7.8 Number of background detected as outliers by all algorithms

```
[ ]: len(df_outliers[(df_outliers['attack'] == False) & (df_outliers['if_outliers']_
↪ == True) & (df_outliers['lof_outliers'] == True) &_
↪ (df_outliers['rc_outliers'] == True)])
```

```
[ ]: 68
```